

freestyle

June 27, 2024

1 Machine Learning for Finance Freestyle

In this lab you'll be given the opportunity to apply everything you have learned to build a trading strategy for SP500 stocks. First, let's introduce the dataset you'll be using.

1.1 The Data

Use BigQuery's magic function to pull data as follows:

Dataset Name: ml4f

Table Name: percent_change_sp500

The following query will pull 10 rows of data from the table:

```
[1]: %%bigquery df
      SELECT
        *
      FROM
        `cloud-training-prod-bucket.ml4f.percent_change_sp500`
      LIMIT
        10
```

Query is running: 0%| |

Downloading: 0%| |

```
[2]: df.head()
```

```
[2]:  symbol      Date  Open  Close  tomorrow_close  tomo_close_m_close  \
0      A  2003-12-17  27.51  27.18             27.82             0.64
1      A  2006-09-27  32.95  32.27             32.98             0.71
2      A  2006-04-18  37.08  37.99             38.61             0.62
3      A  2011-10-19  33.31  32.99             33.73             0.74
4      A  2002-01-22  28.40  27.16             28.00             0.84

      close_MIN_prior_5_days  close_MIN_prior_20_days  close_MIN_prior_260_days  \
0              0.976085              0.976085              0.421266
1              1.006508              0.943291              0.839170
2              0.956304              0.954725              0.536983
3              0.984238              0.891179              0.891179
```

4	1.048601	1.019882	0.694404
---	----------	----------	----------

	close_MAX_prior_5_days	...	close_STDDEV_prior_20_days	\
0	1.010302	...	0.018430	
1	1.025101	...	0.024930	
2	0.973151	...	0.016506	
3	1.023947	...	0.035434	
4	1.121134	...	0.063416	

	close_STDDEV_prior_260_days	\
0	0.168129	
1	0.089782	
2	0.142754	
3	0.184974	
4	0.365963	

	close_values_prior_260	days_on_market	\
0	[16.95, 16.79, 16.4, 17.16, 17.32, 17.18, 17.2...	1025	
1	[33.03, 33.91, 33.65, 34.03, 33.64, 33.68, 33...	1724	
2	[21.66, 21.58, 21.37, 21.14, 21.36, 20.81, 20...	1611	
3	[33.8, 33.75, 33.68, 34.31, 34.23, 34.48, 34.6...	2999	
4	[56.13, 58.25, 55.06, 53.25, 53.0, 55.31, 57.5...	544	

	scaled_change	s_p_scaled_change	normalized_change	\
0	0.023547	0.011798	0.011749	
1	0.022002	0.001713	0.020289	
2	0.016320	0.002027	0.014293	
3	0.022431	0.004554	0.017877	
4	0.030928	0.007925	0.023003	

	company	industry	direction
0	Agilent Technologies Inc	Health Care	UP
1	Agilent Technologies Inc	Health Care	UP
2	Agilent Technologies Inc	Health Care	UP
3	Agilent Technologies Inc	Health Care	UP
4	Agilent Technologies Inc	Health Care	UP

[5 rows x 26 columns]

As you can see, the table contains daily open and close data for SP500 stocks. The table also contains some features that have been generated for you using [navigation functions](#) and [analytic functions](#). Let's dig into the schema a bit more.

```
[10]: %%bigquery
SELECT
    * EXCEPT(is_generated, generation_expression, is_stored, is_updatable)
FROM
```

```

`cloud-training-prod-bucket.ml4f`.INFORMATION_SCHEMA.COLUMNS
WHERE
  table_name = "percent_change_sp500"

```

```

[10]:
      table_catalog table_schema      table_name \
0    cloud-training-prod-bucket      ml4f percent_change_sp500
1    cloud-training-prod-bucket      ml4f percent_change_sp500
2    cloud-training-prod-bucket      ml4f percent_change_sp500
3    cloud-training-prod-bucket      ml4f percent_change_sp500
4    cloud-training-prod-bucket      ml4f percent_change_sp500
5    cloud-training-prod-bucket      ml4f percent_change_sp500
6    cloud-training-prod-bucket      ml4f percent_change_sp500
7    cloud-training-prod-bucket      ml4f percent_change_sp500
8    cloud-training-prod-bucket      ml4f percent_change_sp500
9    cloud-training-prod-bucket      ml4f percent_change_sp500
10   cloud-training-prod-bucket      ml4f percent_change_sp500
11   cloud-training-prod-bucket      ml4f percent_change_sp500
12   cloud-training-prod-bucket      ml4f percent_change_sp500
13   cloud-training-prod-bucket      ml4f percent_change_sp500
14   cloud-training-prod-bucket      ml4f percent_change_sp500
15   cloud-training-prod-bucket      ml4f percent_change_sp500
16   cloud-training-prod-bucket      ml4f percent_change_sp500
17   cloud-training-prod-bucket      ml4f percent_change_sp500
18   cloud-training-prod-bucket      ml4f percent_change_sp500
19   cloud-training-prod-bucket      ml4f percent_change_sp500
20   cloud-training-prod-bucket      ml4f percent_change_sp500
21   cloud-training-prod-bucket      ml4f percent_change_sp500
22   cloud-training-prod-bucket      ml4f percent_change_sp500
23   cloud-training-prod-bucket      ml4f percent_change_sp500
24   cloud-training-prod-bucket      ml4f percent_change_sp500
25   cloud-training-prod-bucket      ml4f percent_change_sp500

```

```

      column_name ordinal_position is_nullable data_type \
0              symbol              1         YES   STRING
1              Date                2         YES    DATE
2              Open                3         YES  FLOAT64
3              Close               4         YES  FLOAT64
4      tomorrow_close              5         YES  FLOAT64
5      tomo_close_m_close           6         YES  FLOAT64
6      close_MIN_prior_5_days       7         YES  FLOAT64
7      close_MIN_prior_20_days      8         YES  FLOAT64
8      close_MIN_prior_260_days     9         YES  FLOAT64
9      close_MAX_prior_5_days       10        YES  FLOAT64
10     close_MAX_prior_20_days      11        YES  FLOAT64
11     close_MAX_prior_260_days     12        YES  FLOAT64
12     close_AVG_prior_5_days       13        YES  FLOAT64
13     close_AVG_prior_20_days      14        YES  FLOAT64

```

14	close_AVG_prior_260_days	15	YES	FLOAT64
15	close_STDDEV_prior_5_days	16	YES	FLOAT64
16	close_STDDEV_prior_20_days	17	YES	FLOAT64
17	close_STDDEV_prior_260_days	18	YES	FLOAT64
18	close_values_prior_260	19	NO	ARRAY<FLOAT64>
19	days_on_market	20	YES	INT64
20	scaled_change	21	YES	FLOAT64
21	s_p_scaled_change	22	YES	FLOAT64
22	normalized_change	23	YES	FLOAT64
23	company	24	YES	STRING
24	industry	25	YES	STRING
25	direction	26	YES	STRING

	is_hidden	is_system_defined	is_partitioning_column	\
0	NO	NO	NO	
1	NO	NO	NO	
2	NO	NO	NO	
3	NO	NO	NO	
4	NO	NO	NO	
5	NO	NO	NO	
6	NO	NO	NO	
7	NO	NO	NO	
8	NO	NO	NO	
9	NO	NO	NO	
10	NO	NO	NO	
11	NO	NO	NO	
12	NO	NO	NO	
13	NO	NO	NO	
14	NO	NO	NO	
15	NO	NO	NO	
16	NO	NO	NO	
17	NO	NO	NO	
18	NO	NO	NO	
19	NO	NO	NO	
20	NO	NO	NO	
21	NO	NO	NO	
22	NO	NO	NO	
23	NO	NO	NO	
24	NO	NO	NO	
25	NO	NO	NO	

	clustering_ordinal_position
0	None
1	None
2	None
3	None
4	None

5	None
6	None
7	None
8	None
9	None
10	None
11	None
12	None
13	None
14	None
15	None
16	None
17	None
18	None
19	None
20	None
21	None
22	None
23	None
24	None
25	None

Most of the features, like `open` and `close` are pretty straightforward. The features generated using analytic functions, such as `close_MIN_prior_5_days` are best described using an example. Let's take the 6 most recent rows of data for IBM and reproduce the `close_MIN_prior_5_days` column.

```
[ ]: %%bigquery
SELECT
  *
FROM
  `cloud-training-prod-bucket.ml4f.percent_change_sp500`
WHERE
  symbol = 'IBM'
ORDER BY
  Date DESC
LIMIT 6
```

For `Date = 2013-02-01` how did we arrive at `close_MIN_prior_5_days = 0.989716`? The minimum close over the past five days was 203.07. This is normalized by the current day's close of 205.18 to get `close_MIN_prior_5_days = 203.07 / 205.18 = 0.989716`. The other features utilizing analytic functions were generated in a similar way. Here are explanations for some of the other features:

- **scaled_change**: `tomo_close_m_close / close`
- **s_p_scaled_change**: This value is calculated the same way as **scaled_change** but for the S&P 500 index.
- **normalized_change**: `scaled_change - s_p_scaled_change` The normalization using the S&P index fund helps ensure that the future price of a stock is not due to larger market effects.

Normalization helps us isolate the factors contributing to the performance of a stock_market.

- **direction:** This is the target variable we're trying to predict. The logic for this variable is as follows:

```
CASE
  WHEN normalized_change < -0.01 THEN 'DOWN'
  WHEN normalized_change > 0.01 THEN 'UP'
  ELSE 'STAY'
END AS direction
```

1.2 Create classification model for direction

In this example, your job is to create a classification model to predict the **direction** of each stock. Be creative! You can do this in any number of ways. For example, you can use BigQuery, Scikit-Learn, or AutoML. Feel free to add additional features, or use time series models.

1.2.1 Establish a Simple Benchmark

One way to assess the performance of a model is to compare it to a simple benchmark. We can do this by seeing what kind of accuracy we would get using the naive strategy of just predicting the majority class. Across the entire dataset, the majority class is 'STAY'. Using the following query we can see how this naive strategy would perform.

```
[9]: %%bigquery
WITH subset as (
  SELECT
    Direction
  FROM
    `cloud-training-prod-bucket.ml4f.percent_change_sp500`
  WHERE
    tomorrow_close IS NOT NULL
)
SELECT
  Direction,
  100.0 * COUNT(*) / (SELECT COUNT(*) FROM subset) as percentage
FROM
  subset
GROUP BY
  Direction
```

```
[9]: Direction  percentage
0      STAY    53.766049
1       UP    23.240681
2     DOWN    22.993271
```

So, the naive strategy of just guessing the majority class would have accuracy of around 54% across the entire dataset. See if you can improve on this.